


# Gaussian Mixture Models

# K-Means Clustering

- Algorithm organizes data into clusters such that there is high intra-cluster similarity and low inter-cluster similarity.
  - A datapoint will belong to one cluster, not several – resulting in a specific number of disjoint non-hierarchical clusters.
- Initialized prototypes, then iterated between two phases:
  - Expectation-step: Each data point assigned to nearest prototype
  - Maximization-step: Prototypes updated to be the cluster means



# Limitations of K-Means


- Clusters are assumed to be spherical and equally sized, which is not valid in most real-world scenarios.
- K-Means is a hard clustering method, with each data point being assigned to a single cluster.
  - A small shift of a data point can flip it to a different cluster.
  - There is no uncertainty measure or probability that tells us how much a data point is associated with a specific cluster.
- Not clear on how to choose the value of K
- Probable solution is to replace the “hard” clustering of K-means with ‘soft’ probabilistic assignments
- The probability distribution of the data is represented as a Gaussian Mixture Model.

# Understanding the Different Approaches

## Scenario 1:

We have labels and we would like to model the appearance of the two classes –

### Maximum Likelihood Estimation



- Assuming observed data points are generated independently,  $D = \{x_1, x_2, \dots, x_N\}$

$$\text{Likelihood} = p(x_1, \dots, x_N; \mu_i, \sigma_i^2) = \prod_{n=1}^N \frac{1}{\sqrt{2\pi}\sigma_i} e^{-\frac{(x_n - \mu_i)^2}{2\sigma_i^2}}$$

where  $i$  is for Class 0 and Class 1 respectively

$$\text{Likelihood} = p(x_1, \dots, x_N; \boldsymbol{\mu}_i, \boldsymbol{\Sigma}_i) = \prod_{n=1}^N \frac{1}{(2\pi)^{\frac{d}{2}} |\boldsymbol{\Sigma}_i|^{\frac{1}{2}}} e^{-\frac{1}{2}(\mathbf{x}_n - \boldsymbol{\mu}_i)^T \boldsymbol{\Sigma}_i^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_i)}$$

# Gaussian Density Function

- We start with the Gaussian Density Function,

$$p(x_1, \dots, x_N; \mu_i, \Sigma_i) = \mathcal{N}(\mathbf{x} | \mu_i, \Sigma_i) = \prod_{n=1}^N \frac{1}{(2\pi)^{\frac{d}{2}} |\Sigma_i|^{\frac{1}{2}}} e^{\left(-\frac{1}{2}(\mathbf{x}-\mu_i)^T \Sigma_i^{-1} (\mathbf{x}-\mu_i)\right)}$$

- $x_n$  represents our data points,
- $d$  is the number of dimensions of each data sample (# of features)
- $\mu$  and  $\Sigma$  are the mean and covariance respectively.
- Say if you have a dataset with 1500 samples and 5 features,  $N = 1500, d = 5$

$$\mathbf{x} = 1500 \times 5, \quad \mu = 1 \times 5, \quad \Sigma = 5 \times 5 \text{ matrices}$$

# Maximum Likelihood

- We set the parameters by maximizing the likelihood function, which we have discussed before is equivalent to maximizing the log likelihood

$$\text{Likelihood} = p(x_1, \dots, x_N; \boldsymbol{\mu}_i, \Sigma_i) = \prod_{n=1}^N \frac{1}{(2\pi)^{\frac{d}{2}} |\Sigma_i|^{\frac{1}{2}}} e^{\left(-\frac{1}{2}(x_n - \boldsymbol{\mu}_i)^T \Sigma_i^{-1} (x_n - \boldsymbol{\mu}_i)\right)}$$

$$\ln p(x_1, \dots, x_N; \boldsymbol{\mu}_i, \Sigma_i) = -\frac{N}{2} \ln |\Sigma_i| - \frac{N}{2} d \ln(2\pi) - \frac{1}{2} \sum_{n=1}^N (x_n - \boldsymbol{\mu}_i)^T \Sigma_i^{-1} (x_n - \boldsymbol{\mu}_i)$$

- Maximizing w.r.t. the mean and the covariance gives the sample mean and sample covariance respectively

$$\boldsymbol{\mu}_i = \frac{1}{N} \sum_{n=1}^N x_n,$$


$$\Sigma_i = \frac{1}{N} \sum_{n=1}^N (x_n - \boldsymbol{\mu}_i) (x_n - \boldsymbol{\mu}_i)^T$$

# Understanding the Different Approaches

## Scenario 2:

Once we have modelled the 2 classes, how do we compute the likelihood that a point belongs to Class 1 or Class 0 –

### Probabilistic Inference



- We compute the likelihood that a particular model generated a sample – Bayes Theorem

$$p(C_0|x_j) = \frac{p(x_j|C_0)p(C_0)}{p(x_j|C_0)p(C_0) + p(x_j|C_1)p(C_1)}, \quad p(C_1|x_j) = \frac{p(x_j|C_1)p(C_1)}{p(x_j|C_0)p(C_0) + p(x_j|C_1)p(C_1)}$$


- Parameters of the Gaussian being known, we can compute the soft assignments of all the points.

# Understanding the Different Approaches

## Scenario 3:

How can we get the labels and the distribution of the classes both at once?

### Expectation Maximization (EM) Algorithm



- We introduce the concept of Gaussian Mixture which is a function comprised of several Gaussians, depending on the number of clusters of our data.
- Each Gaussian is comprised of the following features:
  - Mean  $\mu$  defining the centre of the Gaussian.
  - Covariance  $\Sigma$  defining the width of the Gaussian.
  - An additional mixing parameter  $\pi$  defining how big or small the Gaussian function will be in terms of probability,  $\sum_{k=1}^K \pi_k = 1$ , where  $K$  is the number of clusters.



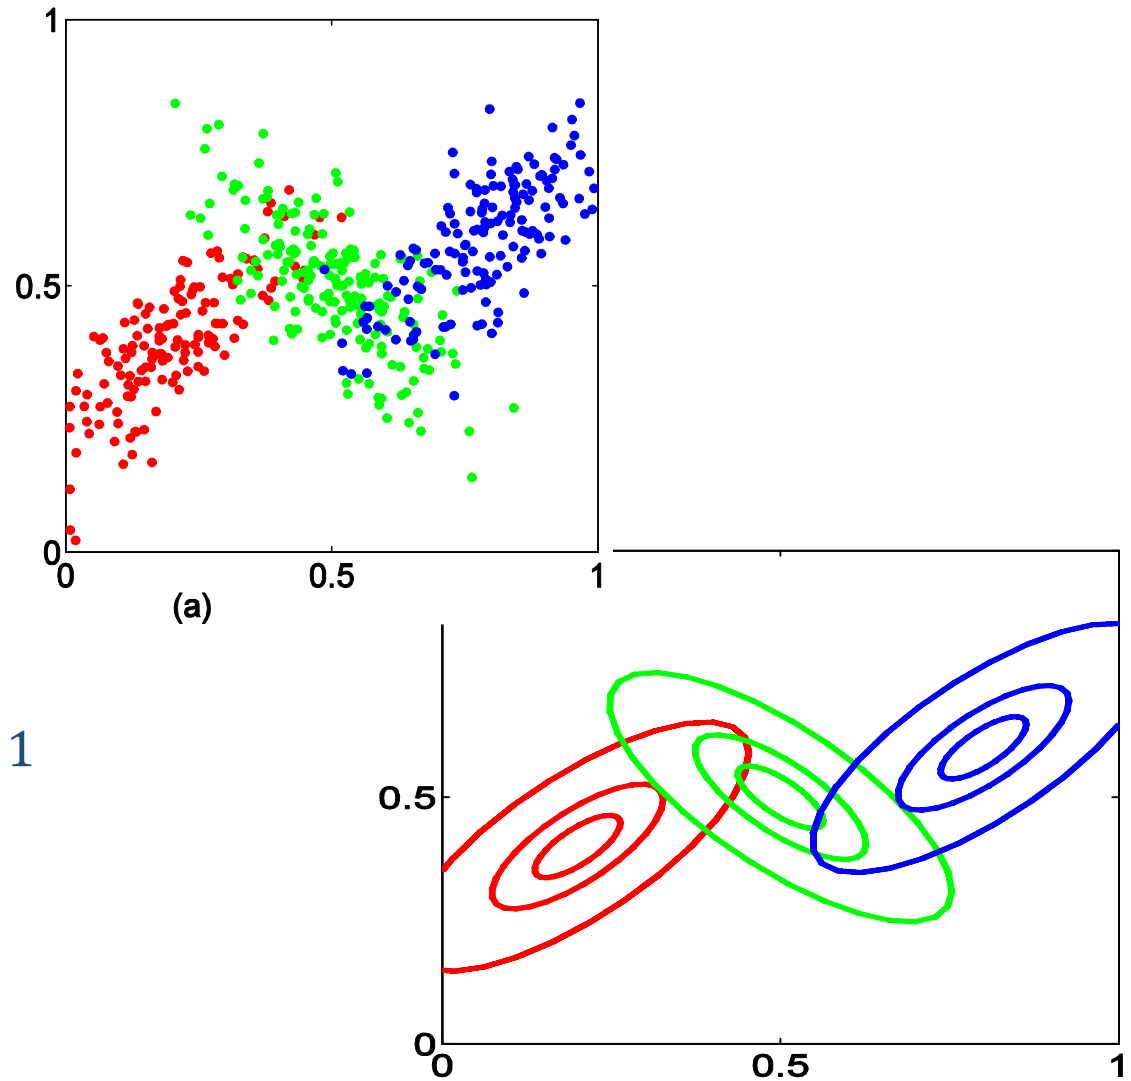
# Gaussian Mixtures

- Linear super-position of Gaussians

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} | \mu_k, \Sigma_k)$$

- Normalization and positivity require

$$\sum_{k=1}^K \pi_k = 1, \quad 0 \leq \pi_k \leq 1$$



# Understanding Mixing Parameter

- The mixing coefficient for a Gaussian is the overall probability of observing a point  $x_n$  that comes from Gaussian  $k$
- Given a data point  $x_n$ , what is the probability it comes from a Gaussian distribution  $k$

$$p(x_n) = \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)$$

- To determine the optimal values for  $\mu_k, \pi_k, \Sigma_k$  we need to determine the maximum likelihood of the model.

$$p(X) = \prod_{n=1}^N p(x_n) = \prod_{n=1}^N \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)$$

# Understanding Mixing Parameter

- Applying the logarithm to both side of the equation, like before:

$$\ln p(X) = \sum_{n=1}^N \ln \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)$$

- Differentiating the above equation with respect to the parameters  $\mu_k, \pi_k, \Sigma_k$  is not straightforward.
- The iterative method, Expectation – Maximization is used to compute the above equation.

# Expectation – Maximization Algorithm

- Step 1: Decide the number of clusters and initialize the mean, covariance, and mixing parameter ( $\mu_k, \Sigma_k, \pi_k$ ) per cluster.
  - The mixing parameter is assumed to be equal for all the clusters in the beginning, say if  $K = 3, \pi = \left\{\frac{1}{3}, \frac{1}{3}, \frac{1}{3}\right\}$ .
- Step 2: **Expectation Step (E step)** – For each sample data  $x_i$ , we calculate the probability that the data point belongs to cluster  $k$ . [[Scenario 2](#)]

$$p(x_{ik}) = \frac{\pi_k \mathcal{N}(x_i | \mu_k, \Sigma_k)}{\sum_{k=1}^K \pi_k \mathcal{N}(x_i | \mu_k, \Sigma_k)}$$

# Expectation – Maximization Algorithm

- Step 2: The E-step computes the probabilities using the current estimates of the model's parameters. It measures how much the  $k^{\text{th}}$  Gaussian distribution is responsible for the generation of the  $i^{\text{th}}$  data point.
- Step 3: **Maximization Step (M step)** – The algorithm uses the measurement obtained for each data point and each Gaussian distribution in the E-step to update the estimates of the model parameters. [**Scenario -1**]

$$\pi_k = \frac{\sum_{i=1}^N p(x_{ik})}{N}, \quad \mu_k = \frac{\sum_{i=1}^N p(x_{ik}) x_i}{\sum_{i=1}^N p(x_{ik})}, \quad \Sigma_k = \frac{\sum_{i=1}^N p(x_{ik}) (x_i - \mu_k)^2}{\sum_{i=1}^N p(x_{ik})}.$$

# Expectation – Maximization Algorithm

- Step 2: Expectation Step (E-step)
- Step 3: Maximization Step (M- step)
- Step 4: The updated estimates  $\mu_k, \Sigma_k, \pi_k$  are used in the next E-step [Step -2] to compute new responsibilities for the data points.
- The process repeats iteratively until convergence, i.e., the model parameters do not change significantly from one iteration to the next.
- Each step increases the log-likelihood of our model

# Flowchart

Decide on the number of distributions

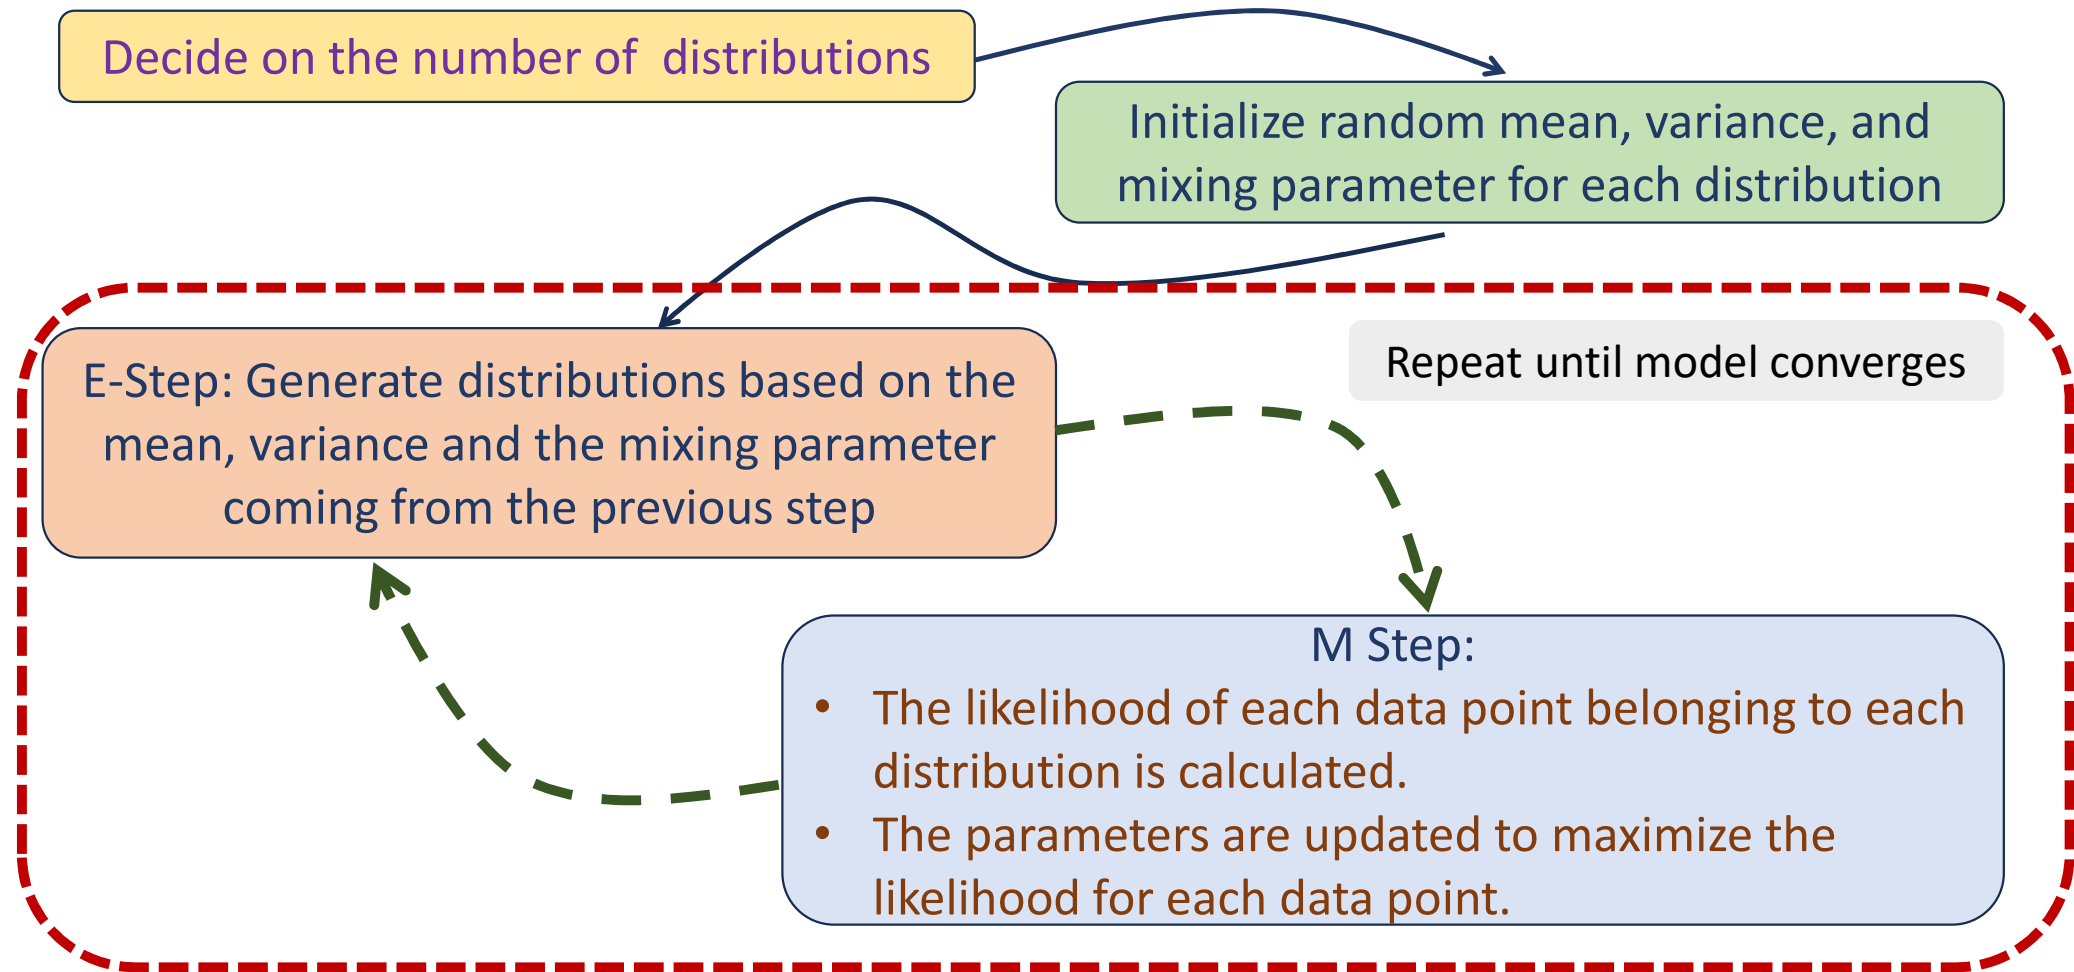
Initialize random mean, variance, and mixing parameter for each distribution

E-Step: Generate distributions based on the mean, variance and the mixing parameter coming from the previous step


Repeat until model converges

M Step:

- The likelihood of each data point belonging to each distribution is calculated.
- The parameters are updated to maximize the likelihood for each data point.




# Gaussian Mixture Example: Start







# After first iteration




# After 2<sup>nd</sup> iteration




After 3<sup>rd</sup> iteration




After 4<sup>th</sup> iteration




After 5<sup>th</sup> iteration



After 6<sup>th</sup> iteration




After 20<sup>th</sup> iteration



# Properties of EM

- EM is trying to optimize a nonconvex function
- But EM is a local optimization algorithm
- Typical solution: Random Restarts
  - Just like K-Means, we run the algorithm many times
  - Each time initialize parameters randomly
  - Pick the parameters that give highest likelihood



# Pros and Cons for EM

## ■ Pros:

- no learning rate (step-size) parameter
- automatically enforces parameter constraints
- very fast for low dimensions
- each iteration guaranteed to improve likelihood

## ■ Cons:

- can get stuck in local minima
- Sensitivity to Initialization
- Computational Complexity
- Choosing the optimal number of clusters