



Training an MLP

Forward and Backward Propagation

Flow graph - Forward propagation



How do we evaluate
our prediction?



$$\begin{aligned} z^{(2)} &= w^{(1)}x \\ a^{(2)} &= f(z^{(2)}) \\ z^{(3)} &= w^{(2)}a^{(2)} \\ \hat{y} &= a^{(3)} = f(z^{(3)}) \end{aligned}$$

Loss Function - Examples



Regression:

- Mean Squared Error: $\frac{1}{N} \sum_j (y_j^{\text{actual}} - y_j^{\text{predicted}})^2$

Classification:

- Cross Entropy Loss: $-(y \log y' + (1 - y) \log(1 - y'))$

Actual Value (y)	Predicted Value (y')	Loss
0	0	0
0	1	∞
1	0	∞
1	1	0




Softmax Activation

- For multi-class classification:
 - We need multiple outputs (1 output per class)
 - We would like to estimate the conditional probability $p(y = c | x)$
- We use the softmax activation function at the output:

$$\mathbf{f}(\mathbf{x}) = \mathbf{softmax}(\mathbf{x}) = \left[\frac{e^{x_1}}{\sum_c e^{x_c}} \quad \dots \quad \frac{e^{x_c}}{\sum_c e^{x_c}} \right]$$


- Normalizes the output



Cross Entropy

$$\text{Loss} = - \sum_{i=1}^n t_i \log p_i, \quad \text{for } n \text{ classes}$$

where t_i is the truth label and p_i is the Softmax probability for the i^{th} class.



$$L_{CE} = - \sum_{i=1} t_i \log p_i$$


$$= -[0 \times \log_2 0.11 + 0 \times \log_2 0.0003 + 0 \times \log_2 0.0418 + 0 \times \log_2 0.057 + 1 \times \log_2 0.8388]$$

$$= -\log_2 0.8388 = 0.2536$$

Training - MLP

- **Step 1** – We initialized the network with random weights.
- **Step 2** – Perform forward propagation and determine \hat{y} .
- **Step 3** – Determine the Loss Function, $|y - \hat{y}|$.
- **Step 4** – Do backward propagation and determine change in weights.
- **Step 5** – Update all weights in all layers.
- **Step 6** – Repeat Steps 2 -5 until convergence.

Forward Propagation



$$z^{(2)} = X W^{(1)} \quad (1)$$

$$a^{(2)} = f(z^{(2)}) \quad (2)$$

$$z^{(3)} = a^{(2)} W^{(2)} \quad (3)$$

$$\hat{y} = f(z^{(3)}) \quad (4)$$

$$J = \frac{1}{2} \sum_i (y_i - \hat{y}_i)^2 \quad (5)$$

$$\begin{aligned} J &= \frac{1}{2} \sum_i (y_i - f(z^{(3)}))^2 = \frac{1}{2} \sum_i (y_i - f(a^{(2)} W^{(2)}))^2 = \frac{1}{2} \sum_i (y_i - f(f(z^{(2)}) W^{(2)}))^2 \\ &= \frac{1}{2} \sum_i (y_i - f(f(X W^{(1)}) W^{(2)}))^2 \end{aligned}$$



MLP - Training

$$J = \frac{1}{2} \sum_i \left(y_i - f(f(X W^{(1)}) W^{(2)}) \right)^2$$

- How does J change, when we change $W^{(1)}, W^{(2)}$?

$$\boxed{\frac{\partial J}{\partial W^{(1)}}, \frac{\partial J}{\partial W^{(2)}}} \Rightarrow \frac{\partial J}{\partial W}$$

- Perform Gradient Descent, where the weights are updated, by computing the gradient of the error function J at W .



MLP:: Gradient Descent


- Stochastic Gradient Descent

$$\begin{bmatrix} 8 & 3 \\ 2 & 11 \\ 9 & 5 \end{bmatrix} \begin{matrix} \rightarrow \frac{\partial J}{\partial W} \rightarrow \text{update } W \\ \rightarrow \frac{\partial J}{\partial W} \rightarrow \text{update } W \\ \rightarrow \frac{\partial J}{\partial W} \rightarrow \text{update } W \end{matrix}$$

- Batch Gradient Descent

$$\begin{bmatrix} 8 & 3 \\ 2 & 11 \\ 9 & 5 \end{bmatrix} \begin{matrix} \rightarrow \frac{\partial J}{\partial W} \\ \rightarrow \frac{\partial J}{\partial W} \\ \rightarrow \frac{\partial J}{\partial W} \end{matrix} + \sum \frac{\partial J}{\partial W} \rightarrow \text{update } W$$

MLP : Backward Propagation




$$J = \frac{1}{2} \sum_i \left(y_i - f(f(X W^{(1)}) W^{(2)}) \right)^2$$

- How do we compute $\frac{\partial J}{\partial W^{(1)}}$, $\frac{\partial J}{\partial W^{(2)}}$?

- J is a function of $W^{(2)}$
- J is a function of a function of $W^{(1)}$

We use Chain Rule to compute these partial derivatives

Backward Propagation



$$z^{(2)} = X W^{(1)} \quad (1)$$

$$a^{(2)} = f(z^{(2)}) \quad (2)$$

$$z^{(3)} = a^{(2)} W^{(2)} \quad (3)$$

$$\hat{y} = f(z^{(3)}) \quad (4)$$

$$J = \frac{1}{2} \sum_i (y_i - \hat{y}_i)^2 \quad (5)$$

$$J = \frac{1}{2} \sum_i \left(y_i - f(f(X W^{(1)}) W^{(2)}) \right)^2$$

$$\frac{\partial J}{\partial W^{(2)}} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}} \frac{\partial z^{(3)}}{\partial W^{(2)}}$$

$$\frac{\partial J}{\partial W^{(1)}} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}} \frac{\partial z^{(3)}}{\partial a^{(2)}} \frac{\partial a^{(2)}}{\partial z^{(2)}} \frac{\partial z^{(2)}}{\partial W^{(1)}}$$

Backward Propagation

$$\frac{\partial J}{\partial W^{(2)}} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}} \frac{\partial z^{(3)}}{\partial W^{(2)}}$$

$$= \sum_i (y - \hat{y}) f'(z^{(3)}) a^{(2)}$$

$$\frac{\partial J}{\partial W^{(1)}} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}} \frac{\partial z^{(3)}}{\partial a^{(2)}} \frac{\partial a^{(2)}}{\partial z^{(2)}} \frac{\partial z^{(2)}}{\partial W^{(1)}}$$

$$= \sum_i (y - \hat{y}) f'(z^{(3)}) W^{(2)} f'(z^{(2)}) X$$

$$z^{(2)} = X W^{(1)} \quad (1)$$

$$a^{(2)} = f(z^{(2)}) \quad (2)$$

$$z^{(3)} = a^{(2)} W^{(2)} \quad (3)$$

$$\hat{y} = f(z^{(3)}) \quad (4)$$

$$J = \frac{1}{2} \sum_i (y_i - \hat{y}_i)^2 \quad (5)$$

Training - MLP

- **Step 1** – We initialized the network with random weights.

- **Step 2** – Forward Propagation

$$z^{(2)} = X W^{(1)} \quad (1)$$

$$a^{(2)} = f(z^{(2)}) \quad (2)$$

$$z^{(3)} = a^{(2)} W^{(2)} \quad (3)$$

$$\hat{y} = f(z^{(3)}) \quad (4)$$

$$J = \frac{1}{2} \sum_i (y_i - \hat{y}_i)^2 \quad (5)$$

- **Step 3** – Backward Propagation

$$\frac{\partial J}{\partial W^{(2)}} = \sum_i (y - \hat{y}) f'(z^{(3)}) a^{(2)}$$

$$\frac{\partial J}{\partial W^{(1)}} = \sum_i (y - \hat{y}) f'(z^{(3)}) W^{(2)} f'(z^{(2)}) X$$



- **Step 4** – Update all weights simultaneously

$$w^{(t+1)} = w^{(t)} - \eta^{(t)} \nabla_w J(w)$$

- Repeat Step – 2, 3 and 4 until convergence

Vanishing/Exploding Gradient Problem

- Backpropagated errors multiply at each layer, resulting in exponential decay (if derivative is small) or growth (if derivative is large).
- Makes it very difficult to train deep networks, or simple recurrent networks over many time steps.



Overfitting in Neural Networks

- Deep Neural Networks are prone to overfitting because of the large number of parameters it encloses.
- The network can become less prone to overfitting by removing certain layers of the network or decreasing the number of neurons.
- Different strategies have been proposed to take care of the overfitting.
 - Reduce overfitting by training the network on more examples.
 - Reduce overfitting by changing the complexity of the network.
 - Changing the network structure (number of weights)
 - Changing the network parameters (values of weights)

Data Augmentation (Jittering)

- Create *virtual* training samples

- Horizontal flip
- Random crop
- Color casting
- Geometric distortion



Data Augmentation



Original Image



De-texturized



De-colored



Edge Enhanced



Salient Edge Map



Flip/Rotate

- Data augmentation is used to artificially increase the size of a training dataset by applying various transformations (flipping, rotation, scaling, cropping, shearing, and adding noise) to the existing data.
- It helps the model to learn invariant features and improves its ability to handle different variations and deformations

Deep Image [Wu et al. 2015]

<http://arxiv.org/pdf/1501.02876v2.pdf>


Weight Regularization

- Penalize the model during training based on the magnitude of the weights.
- Ridge Regression tries to reduce the length $\|w\|$ of the parameter vector, promoting lesser dependency of \hat{y} on predictors (lower model complexity).

$$L_{\text{Ridge}} = J + \alpha \frac{1}{2} \sum_{i=1}^P w_i^2$$

- Lasso Regression tries to reduce the city block distance $|w|$ of the parameter vector, causing automatic feature selection.

$$L_{\text{Lasso}} = J + \alpha \sum_{i=1}^P |w_i|$$




Dropout

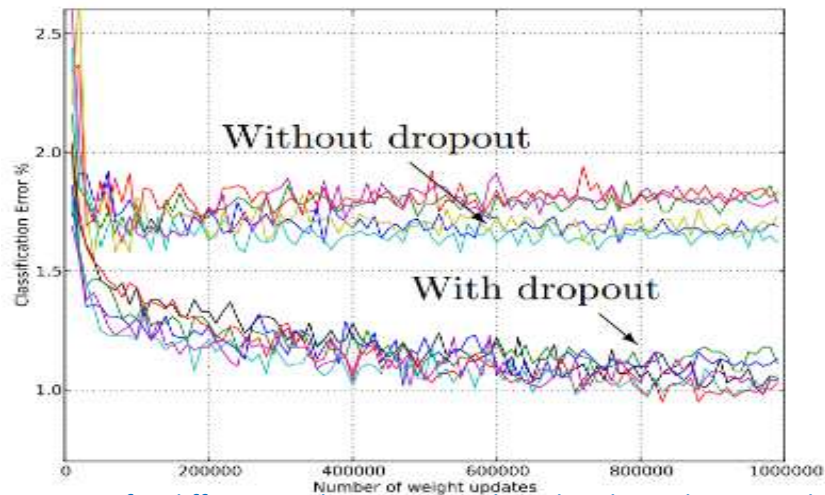

- Used for preventing overfitting in deep neural network which contains a large number of parameters
- Ridge and Lasso reduces overfitting by modifying the loss function, whereas in dropout a certain number of neurons at a layer is deactivated from firing during training
- The dropout rate is a hyperparameter that determines the probability of dropping out each neuron, usually ranging from 0.2 to 0.5.
- Higher dropout rates provide more regularization but may also slow down the convergence of the network.

Dropout: A simple way to prevent neural networks from overfitting [Srivastava JMLR 2014]

<http://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>




Dropout



Test error for different architectures with and without dropout. The networks have 2 to 4 hidden layers each with 1024 to 2048 units.

Main Idea: approximately combining exponentially many different neural network architectures efficiently



Dropout: A simple way to prevent neural networks from overfitting
[[Srivastava JMLR 2014](#)]

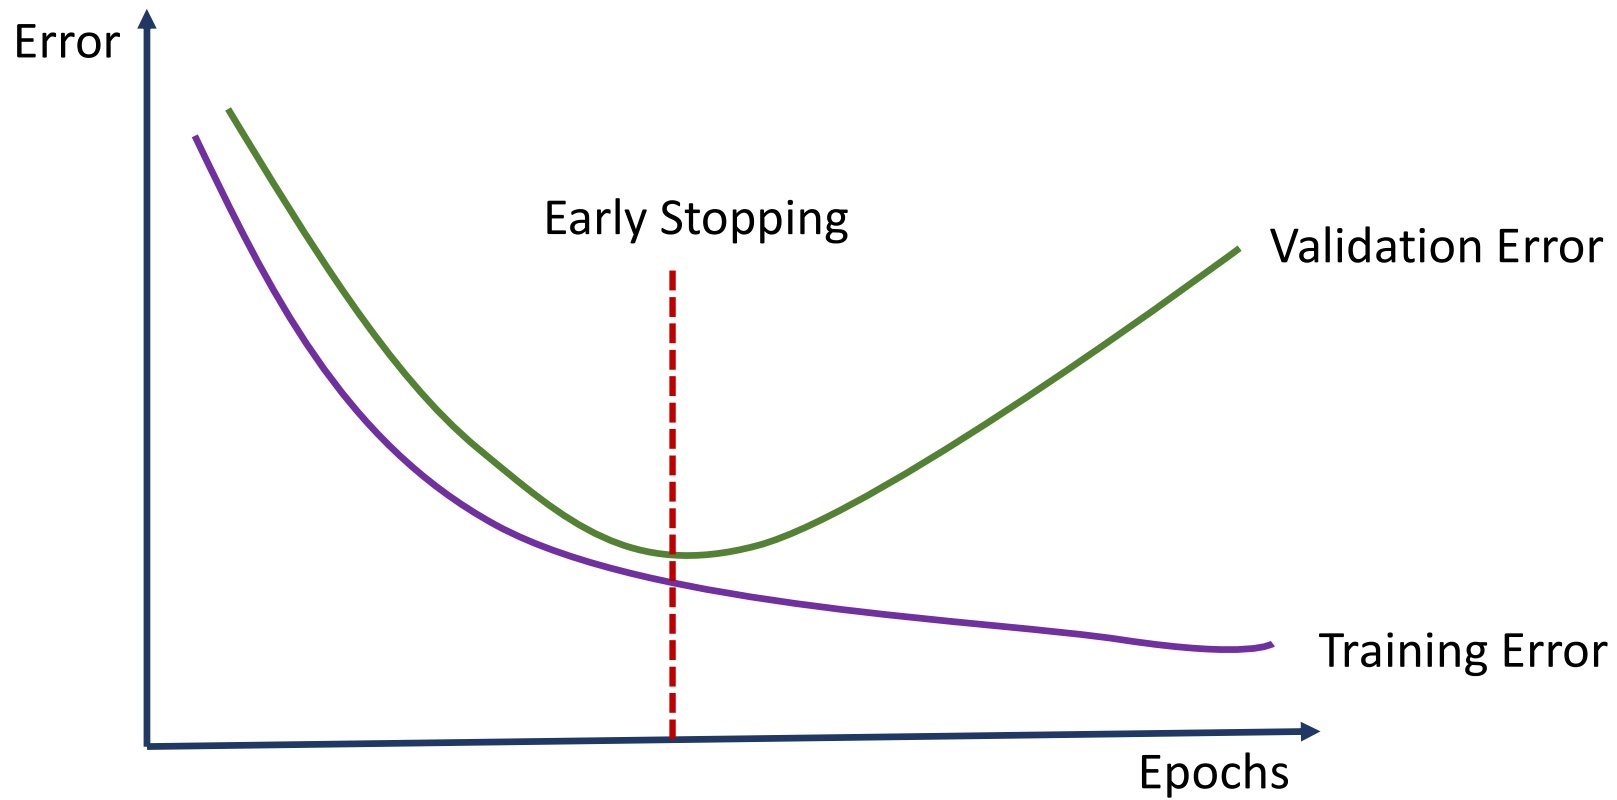
Early Stopping

- While training a large network, at some point the model stops generalizing and starts learning the statistical noise in the training dataset.

Early Stopping

- During training, evaluate the model on a validation dataset after each epoch.
 - Provided the performance of the model on the validation dataset starts to degrade stop the training process instantly.
 - The model at the time the training stopped is used for the test analysis
-
- We train the network on a larger number of training epochs than may normally be required, to give the network plenty of opportunity to fit, then begin to overfit the training dataset.
 - A trigger for stopping the training process is chosen.
 - No change in metric over a given number of epochs.
 - An absolute change in metric or a decrease in performance observed over a certain number of epochs.

Early Stopping




Batch Normalization

- Batch Normalization (BN) is a normalization method/layer for neural networks.
- Usually inputs to neural networks are normalized
- A new layer is added so the gradient can “see” the normalization and made adjustments if needed.
 - The new layer has the power to learn the identity function to de-normalize the features if necessary!

<https://arxiv.org/abs/1502.03167>

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;
Parameters to be learned: γ, β
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$



Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift [[Ioffe and Szegedy 2015](#)]