


# Time-Series Models

Recurrent Neural Network

# Examples

BSE SENSEX



Year	Population(in Million)
1921	251
1931	279
1941	319
1951	361
1961	439
1971	548
1981	685

# Time Series

- Time series is a sequence of observations often ordered in time.
- Popular Problem: Given a sequence, predict future samples.
- Applications:
  - Meteorology,
  - Finance,
  - Marketing etc.
- We want a machine learning model to understand sequences, not samples.
- Assume we have a sequence of measurements, and we want to take  $N$  sequential measurements and predict the next one.

# Notation and Problem

- Notation:  $x[0], x[1], x[2], \dots, x[N]$ .
- $x[t]$ , Where  $t$  is the time or index in the sequence.
- Assumption: Measurement at time  $t$  depends on three previous ones.
  - i.e.,  $t-1, t-2$  and  $t-3$
- Why 3? We can have a different number.

Feature Vector	
Feature	$Y_i$
$V_1$	$X_4$
$V_2$	$X_5$
$V_3$	$X_6$
$V_4$	$X_7$



Rearranged Data			
Feature-1	Feature-2	Feature-3	$Y_i$
$X_1$	$X_2$	$X_3$	$X_4$
$X_2$	$X_3$	$X_4$	$X_5$
$X_3$	$X_4$	$X_5$	$X_6$
$X_4$	$X_5$	$X_6$	$X_7$



Raw Data	
Time	Sample
1	$X_1$
2	$X_2$
3	$X_3$
4	$X_4$
5	$X_5$
6	$X_6$
7	$X_7$

# A Simple Model

- $X[t] = w_1 X[t-1] + w_2 X[t-2] + w_3 X[t-3] + n$

$n$  is noise

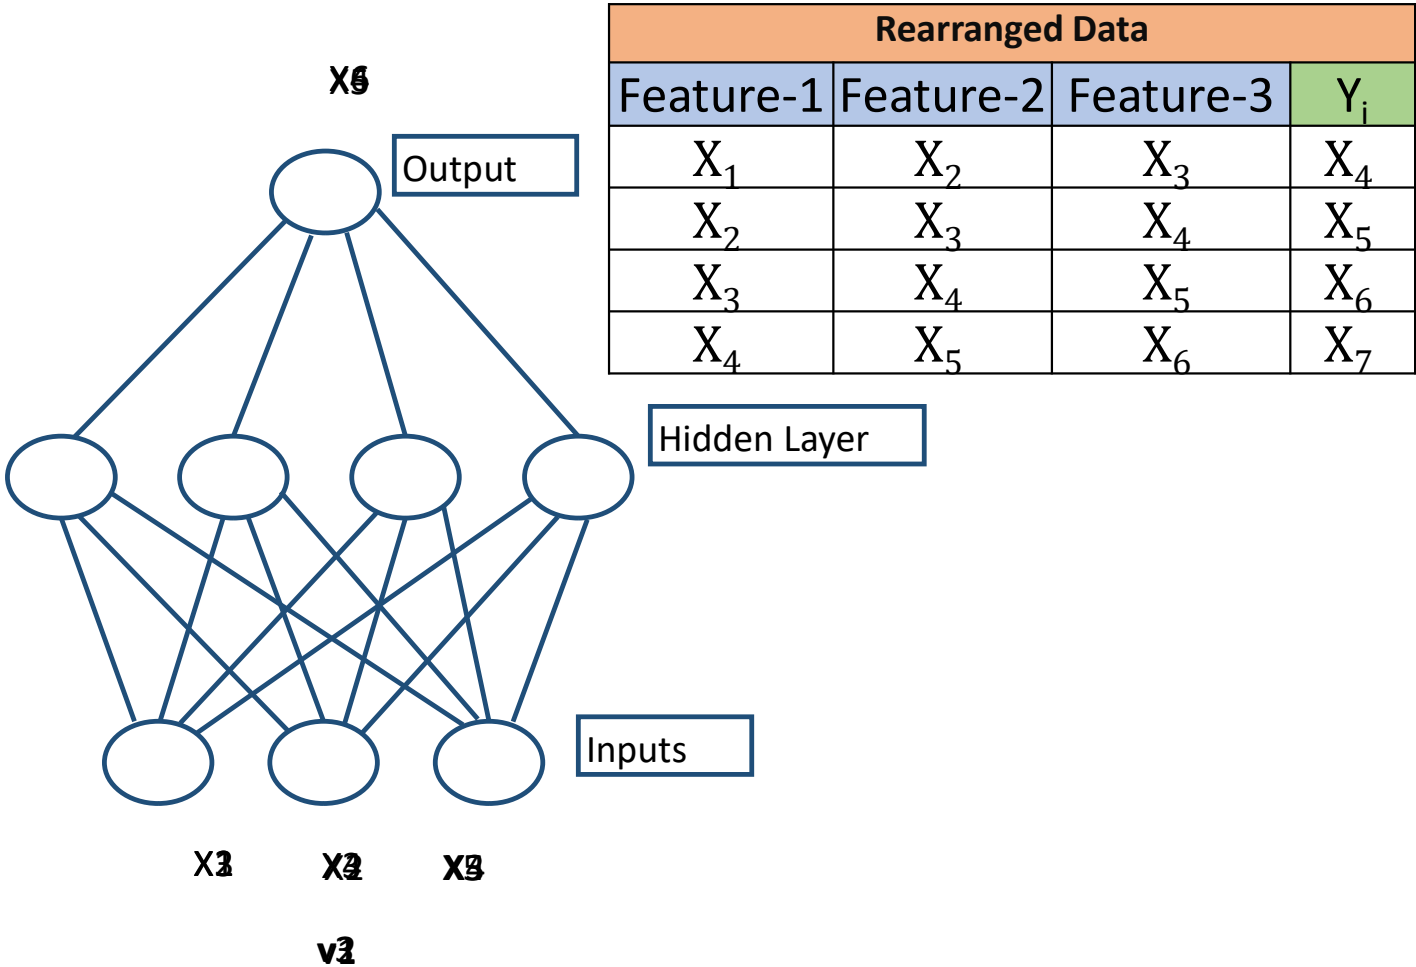
- Given the sequence  $X[0], X[1], \dots, X[N]$ , we find the coefficients  $w_1, w_2, w_3$  such that the prediction error is minimal.

- $X_t = f(W, X_{t-1}, X_{t-2}, X_{t-3}) + n$

$$\Rightarrow \min_W \sum_{t=3}^N (X_t - f(W, X_{t-1}, X_{t-2}, X_{t-3}))^2$$

- Mean Absolute Deviation:  $\sum_{i=1}^N \frac{|X_i - \hat{X}_i|}{N}$
- Mean Absolute Percent Error:  $\frac{100}{N} \sum_{i=1}^N \frac{|X_i - \hat{X}_i|}{\hat{X}_i}$
- Mean Square Error:  $\sum_{i=1}^N \frac{|X_i - \hat{X}_i|^2}{N}$
- Root Mean Square Error:  $\sqrt{\sum_{i=1}^N \frac{|X_i - \hat{X}_i|^2}{N}}$

# Neural Networks for Time Series Forecasting



# Classical Models (AR and MA)

- Auto Regressive (**AR**) Model assumes:  $X_t = \alpha X_{t-1} + \epsilon_t$ , ( $\epsilon_t$  is random uncorrelated)
  - Predict the next term in a sequence from a fixed number of previous terms.
- **AR**: A model of order  $p$  is  $X_t = \sum_{i=1}^p \alpha_i X_{t-i} + \epsilon_t$
- Moving Average (**MA**) model assumes:  $X_t = \epsilon_t + \beta \epsilon_{t-1}$
- **MA**: A model of order  $q$  is  $X_t = \epsilon_t + \sum_{j=1}^q \beta_j \epsilon_{t-j}$

# Classical Models (ARMA & ARIMA)

▪ **ARMA (p,q):**  $X_t = \sum_{i=1}^p \alpha_i X_{t-i} + \sum_{j=0}^q \beta_j \epsilon_{t-j}$ , with  $\beta_0 = 1$

- ARMA is combined from the AR and MA models to model stationary nonseasonal time series data.

▪ **ARIMA (p, d, q):**

- ARIMA is quite similar to ARMA model, with the **I** standing for **Integrated**, i.e. differencing.
- A process is ARIMA (p, q, d) if  $\nabla^d X$  is ARMA (p,q), where  $\nabla X_t = X_t - X_{t-1}$  and  $\nabla^2 X_t = \nabla(\nabla X_t)$
- ARIMA is a combination of a number of differences already applied on the model to make it stationary, the number of previous lags along with residuals errors in order to forecast future values.




# Many Comparisons


- MLP vs ARMA/ARIMA:
  - “Forecasting with artificial neural networks: The state of the art ” – 1998
    - Shows that ANNs are at par or better.
  - “Time series forecasting using a hybrid ARIMA and neural network model” G.P. Zhang (2003)
    - Shows how to get advantages of “both” worlds
- We now know more NN than what we did in 1998 or 2003!!

# Prediction using ARIMA and MLP

## ARIMA



## MLP





# CNNs or MLPs shortcomings


- MLPs/CNNs require fixed input and output size.
- MLPs/CNNs can't classify inputs in multiple places.
- A fully connected network will not distinguish the order and therefore will be missing some information.
- Predicting the next term in a sequence blurs the distinction between supervised and unsupervised learning.
  - Uses method designed for supervised learning, but it doesn't require a separate teaching signal.
  - The network needs to have a memory.

# Memory

- Somehow the computational unit should remember what it has seen before
- We'll call the information the unit's state

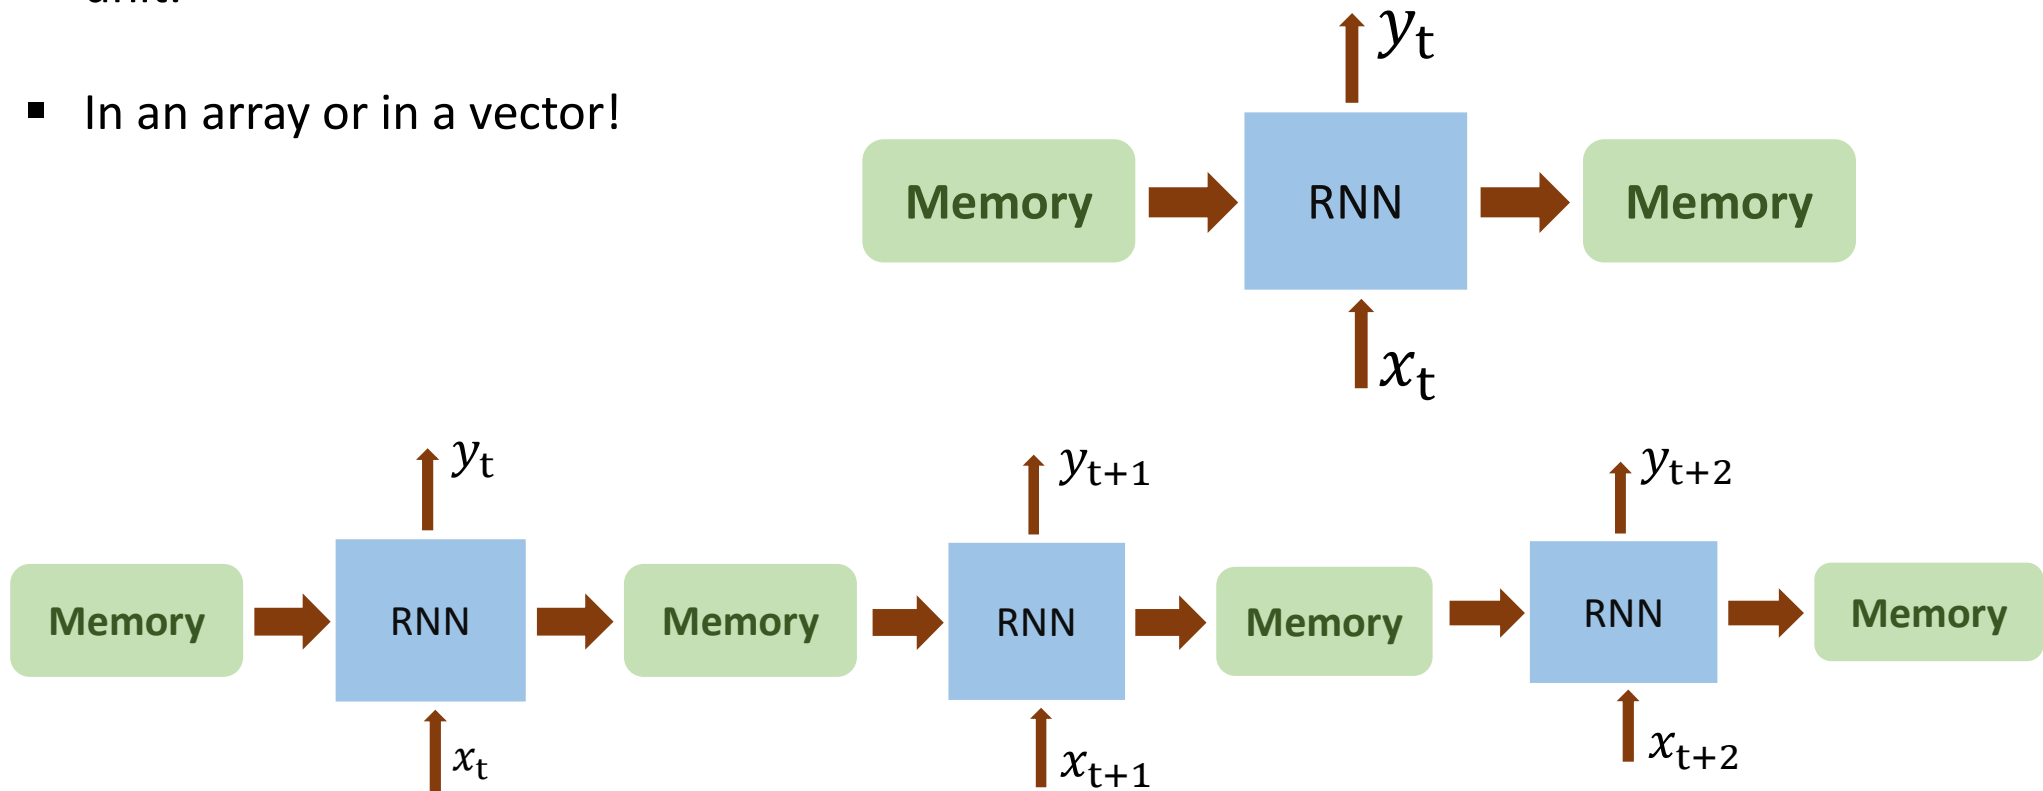


# Handling Individual Time Steps

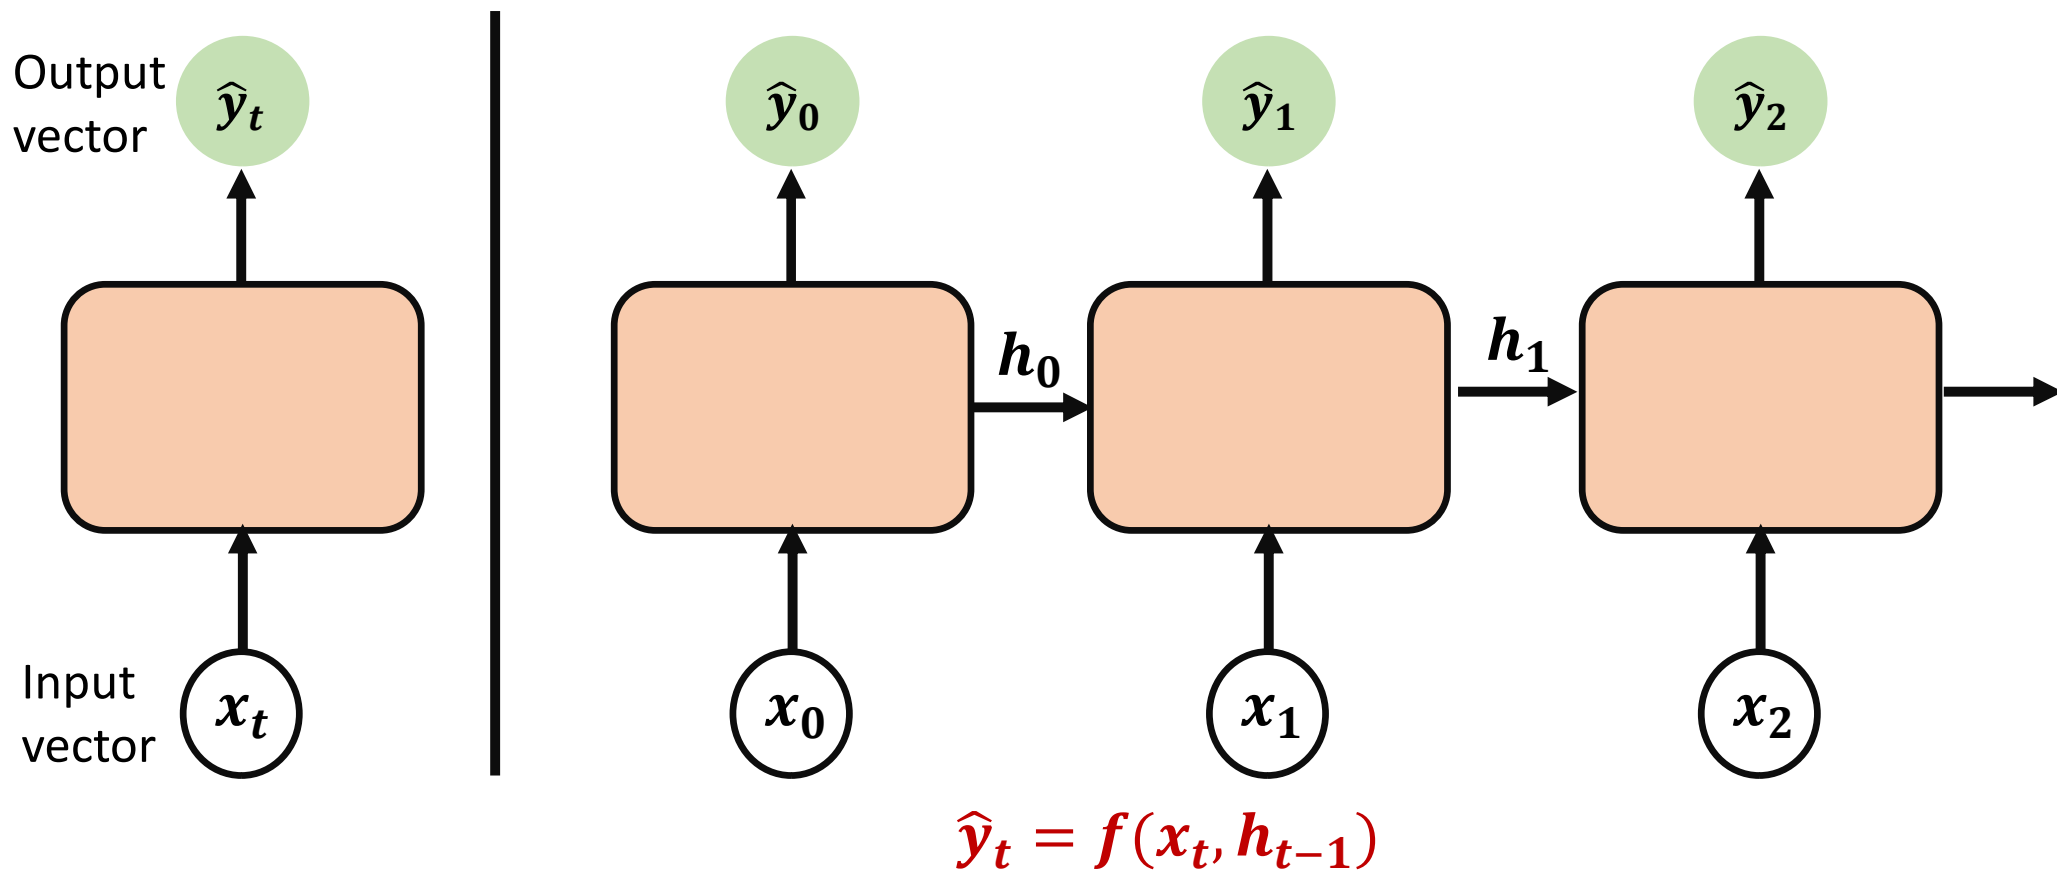


# Recurrent Neural Networks

- The memory or state can be written to a file but in RNNs, we keep it inside the recurrent unit.
- In an array or in a vector!




# Handling Individual Time Steps



# Unrolled RNNs

**Key Idea:** RNNs have an “internal state” that is updated as a sequence is processed

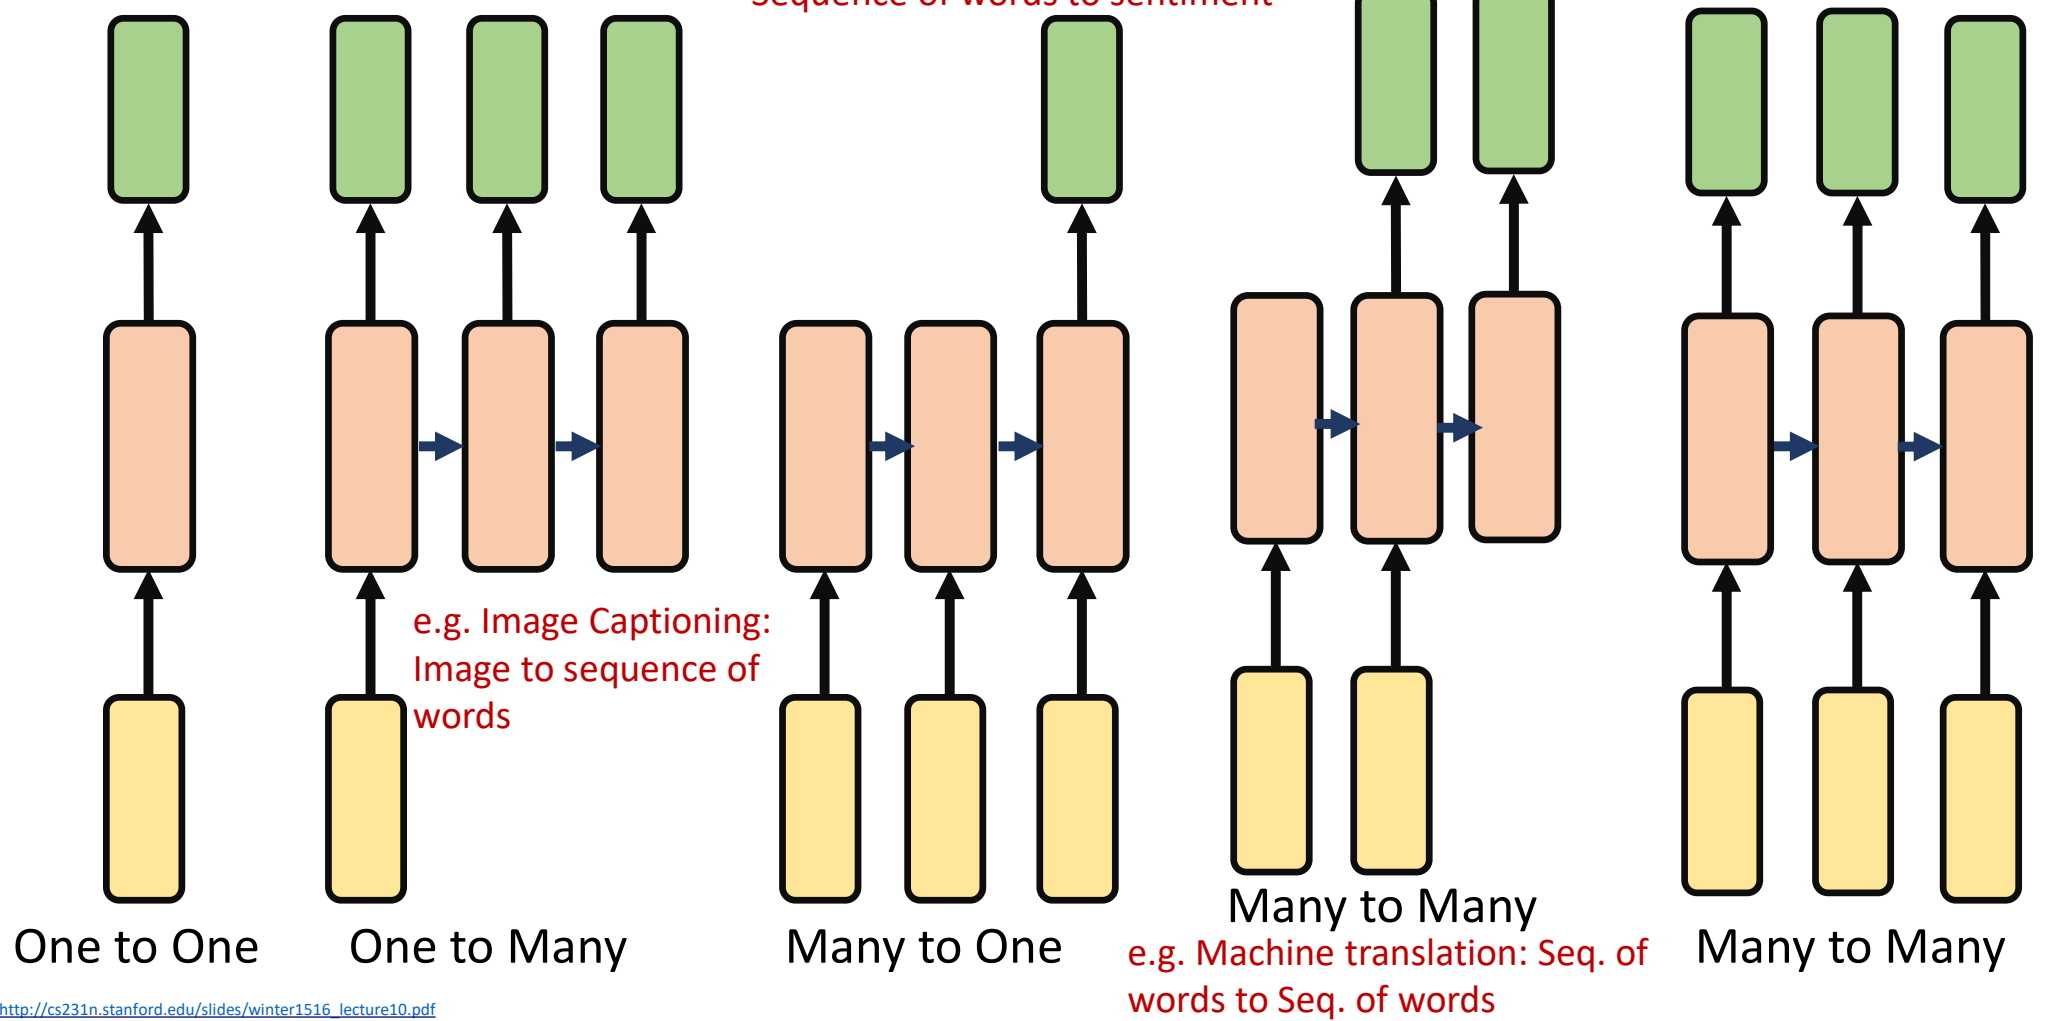
- Temporal dependencies
- Variable Sequence Length



$$\hat{y}_t = f(x_t, h_{t-1})$$



# Modelling a RNN



# RNN hidden state update


- We can process a sequence of vectors  $\mathbf{x}$  by applying a recurrence formula at every time step:  $\mathbf{h}_t = f_W(\mathbf{h}_{t-1}, \mathbf{x}_t)$

$\mathbf{h}_t$ : new state,  $\mathbf{h}_{t-1}$ : old state,  $\mathbf{x}_t$ : input vector at the time step


$f_W$ : some function with parameters  $W$

Note: The same function and the same set of parameters are used at every time step.

- The output  $y_t$  is represented by another function of parameters,  $W_{hy}$ , where  $\mathbf{y}_t = f_{W_{hy}}(\mathbf{h}_t)$



# RNN State Update and Output



**Output Vector**

$$\hat{y}_t = W_{hy}^T h_t$$

**Update Hidden State**

$$h_t = \tanh(W_{hh}^T h_{t-1} + W_{xh}^T x_t)$$


$$h_t = f_W(h_{t-1}, x_t)$$

**Input Vector**

$x_t$



# Recurrent Neural Network




- 3 sets of parameters -  $W_I$ ,  $W_y$ ,  $W_R$  shared for each time-step.
- Reuse the same weight matrix at every time-step.

# Sequence Modelling: Design Criteria

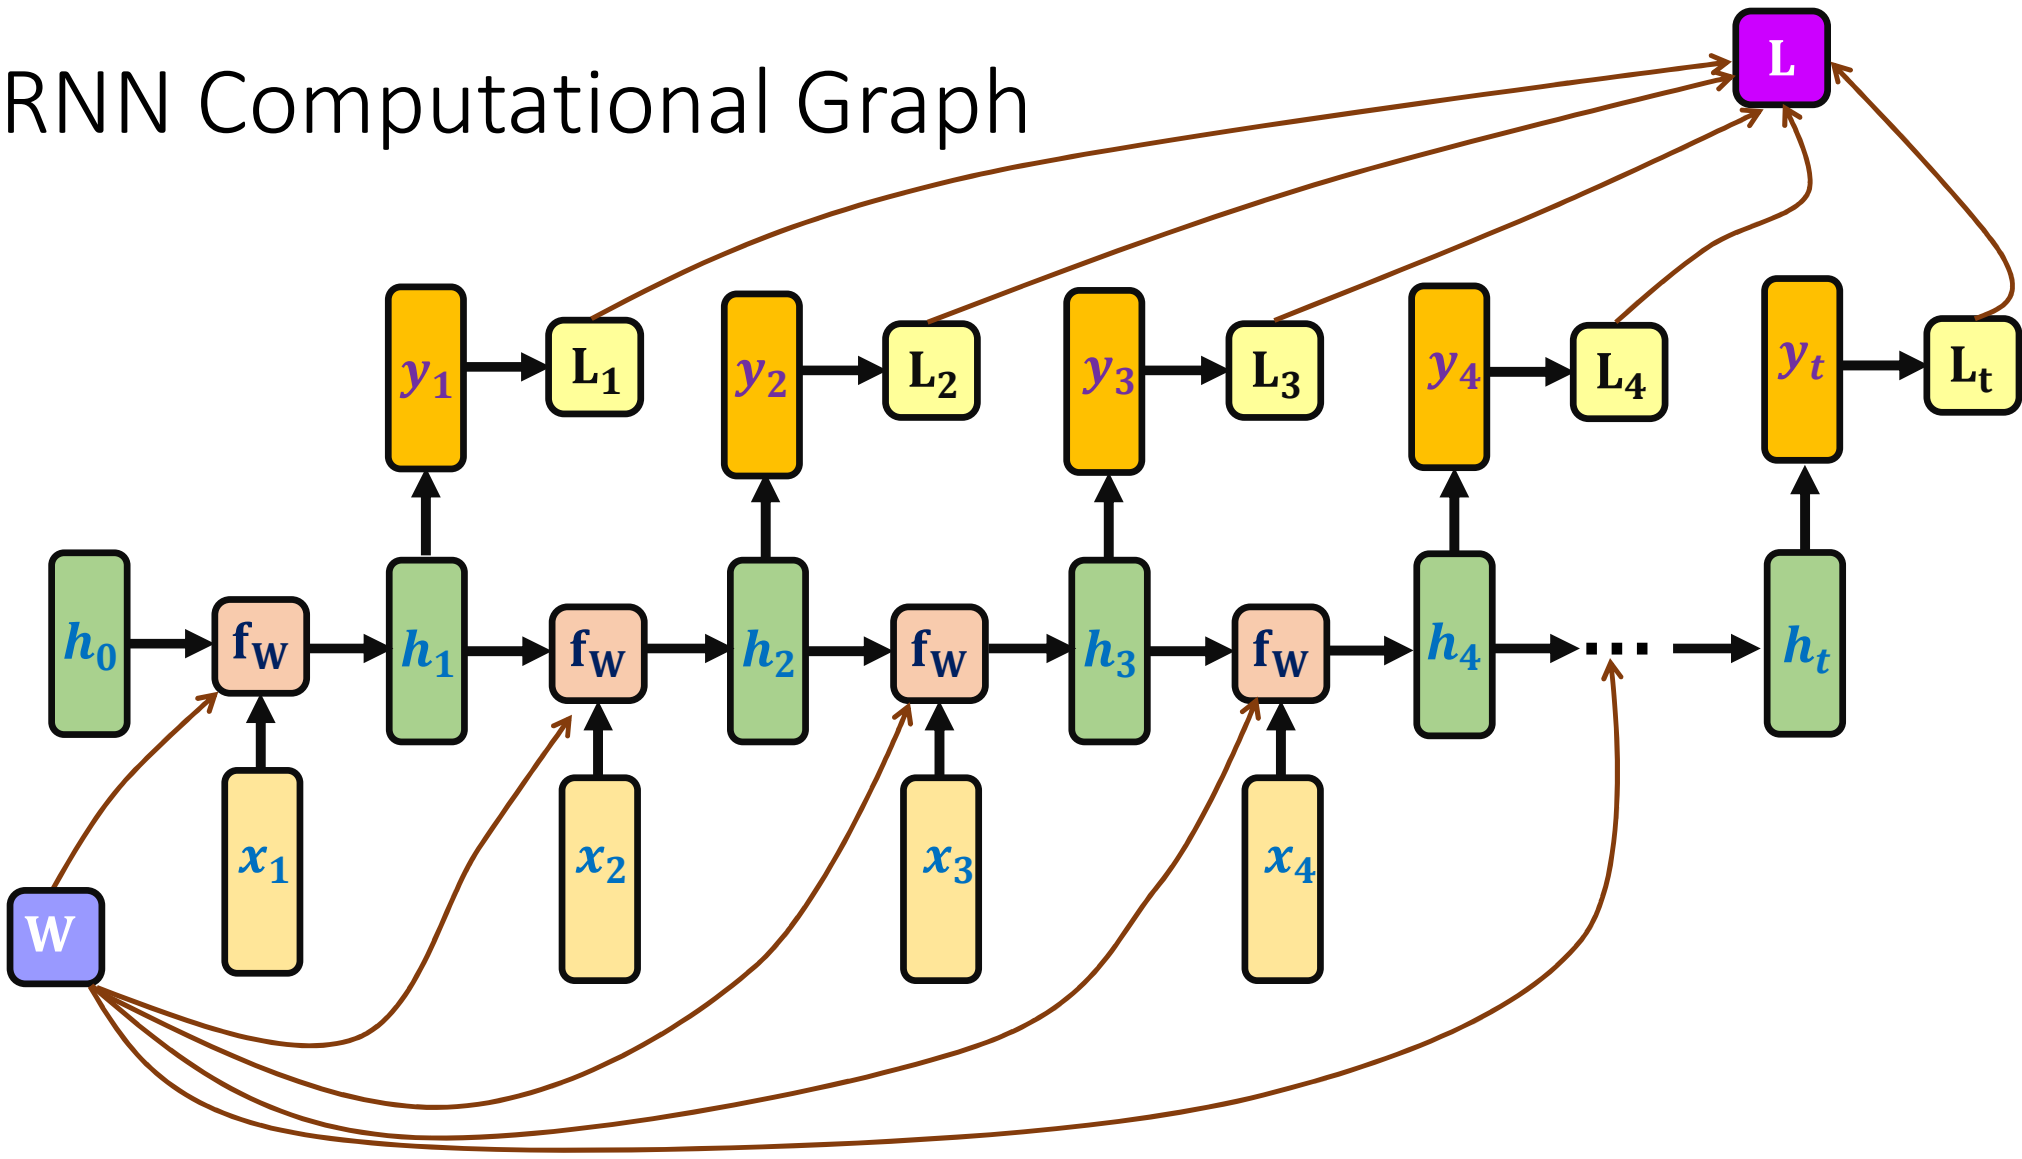
To model sequences, we need to:

1. Handle variable-length sequences
2. Track long-term dependencies
3. Maintain information about order
4. Share parameters across the sequence



Recurrent Neural Networks meets the Sequence Modelling Design Criteria

# RNN Computational Graph




# RNN Gradient Flow

$$\begin{aligned}
 h_t &= \tanh(W_{hh}h_{t-1} + W_{xh}x_t) \\
 &= \tanh\left((W_{hh} \ W_{xh}) \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right) \\
 &= \tanh\left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right)
 \end{aligned}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \tanh'(W_{hh}h_{t-1} + W_{xh}x_t)W_{hh}$$

Backpropagation in time:  $\frac{\partial L}{\partial W} = \sum_{t=1}^T \frac{\partial L_t}{\partial W}$

$$\frac{\partial L_t}{\partial W} = \frac{\partial L_T}{\partial h_T} \frac{\partial h_t}{\partial h_{t-1}} \cdots \frac{\partial h_1}{\partial W} = \frac{\partial L_T}{\partial h_T} \left( \prod_{t=2}^T \frac{\partial h_t}{\partial h_{t-1}} \right) \frac{\partial h_1}{\partial W}$$



# RNN Gradient Flow


$$h_t = \tanh\left(W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}\right)$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \tanh'(W_{hh}h_{t-1} + W_{xh}x_t)W_{hh}$$

Backpropagation in time:

$$\frac{\partial L_t}{\partial W} = \frac{\partial L_T}{\partial h_T} \left( \prod_{t=2}^T \frac{\partial h_t}{\partial h_{t-1}} \right) \frac{\partial h_1}{\partial W} = \frac{\partial L_T}{\partial h_T} \left( \prod_{t=2}^T \boxed{\tanh'(W_{hh}h_{t-1} + W_{xh}x_t)W_{hh}} \right) \frac{\partial h_1}{\partial W}$$

Value almost always less than one,  
vanishing gradient problem





# RNN Gradient Flow

- What if we assumed no non-linearity?

$$h_t = W_{hh}h_{t-1} + W_{xh}x_t$$


$$\frac{\partial h_t}{\partial h_{t-1}} = W_{hh}$$

Backpropagation in time:  $\frac{\partial L}{\partial W} = \sum_{t=1}^T \frac{\partial L_t}{\partial W}$

$$\frac{\partial L_t}{\partial W} = \frac{\partial L_T}{\partial h_T} \left( \prod_{t=2}^T \frac{\partial h_t}{\partial h_{t-1}} \right) \frac{\partial h_1}{\partial W} = \frac{\partial L_T}{\partial h_T} \left( \prod_{t=2}^T W_{hh} \right) \frac{\partial h_1}{\partial W} = \frac{\partial L_T}{\partial h_T} W_{hh}^{T-1} \frac{\partial h_1}{\partial W}$$

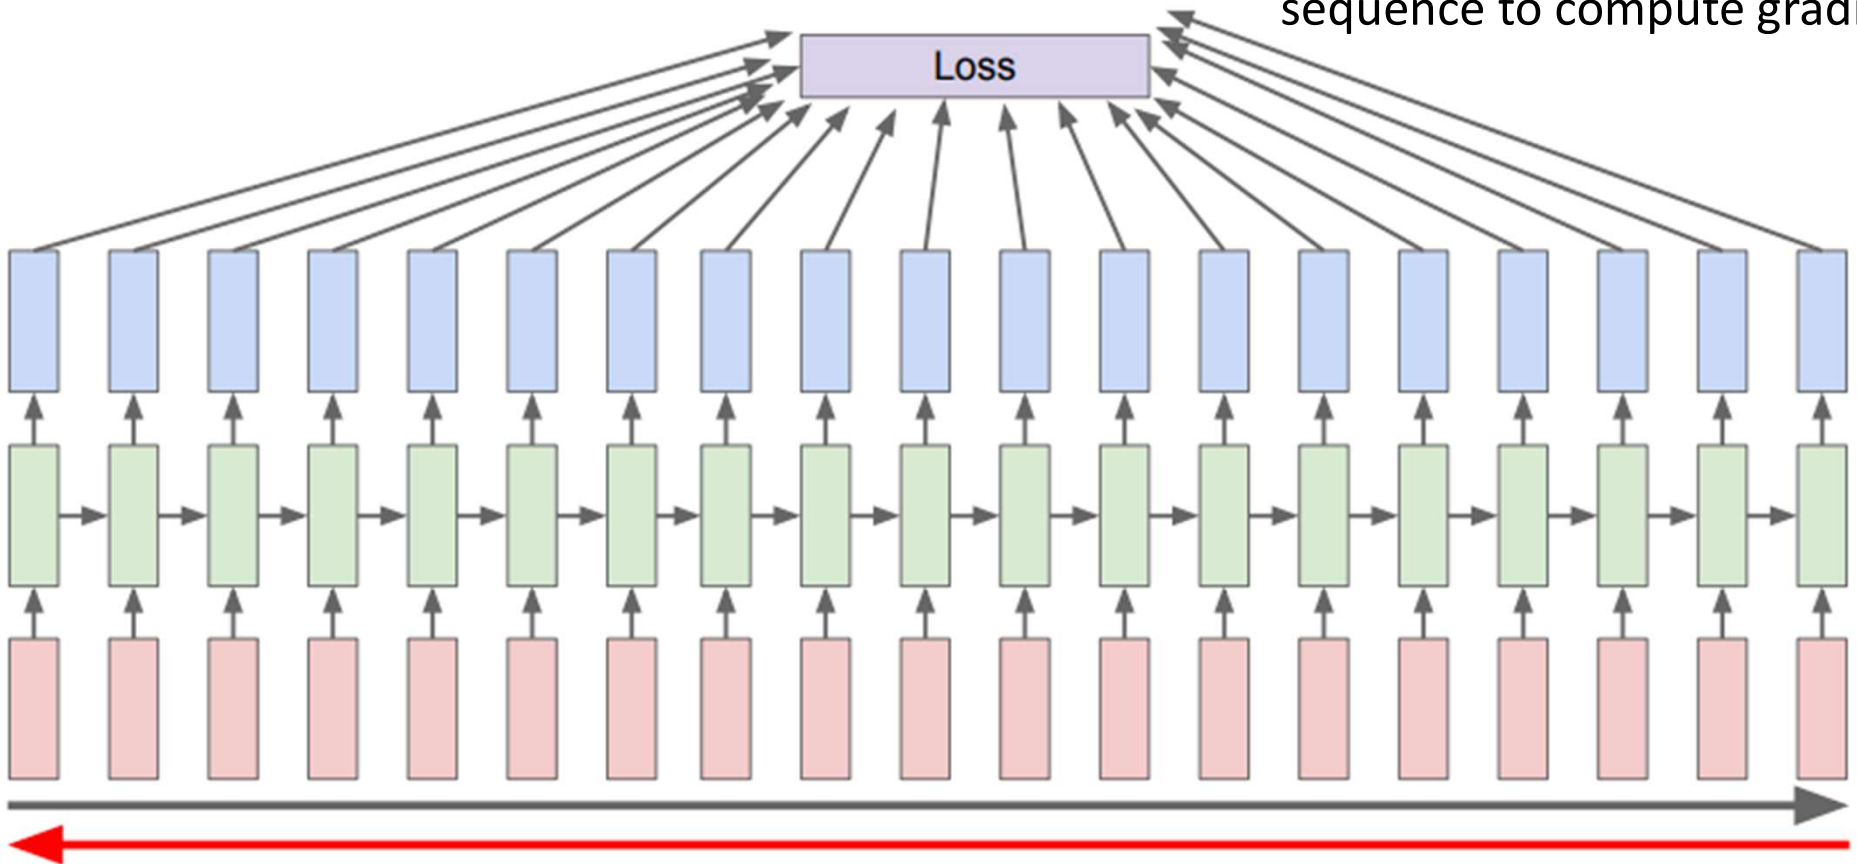
- Largest singular value > 1: Exploding Gradient
- Largest singular value < 1: Vanishing Gradient

Go with gradient clipping, scale gradient if it's norm is too big  
Change RNN architecture




# Backpropagation through time

Forward through entire sequence to compute loss, then backward through entire sequence to compute gradient

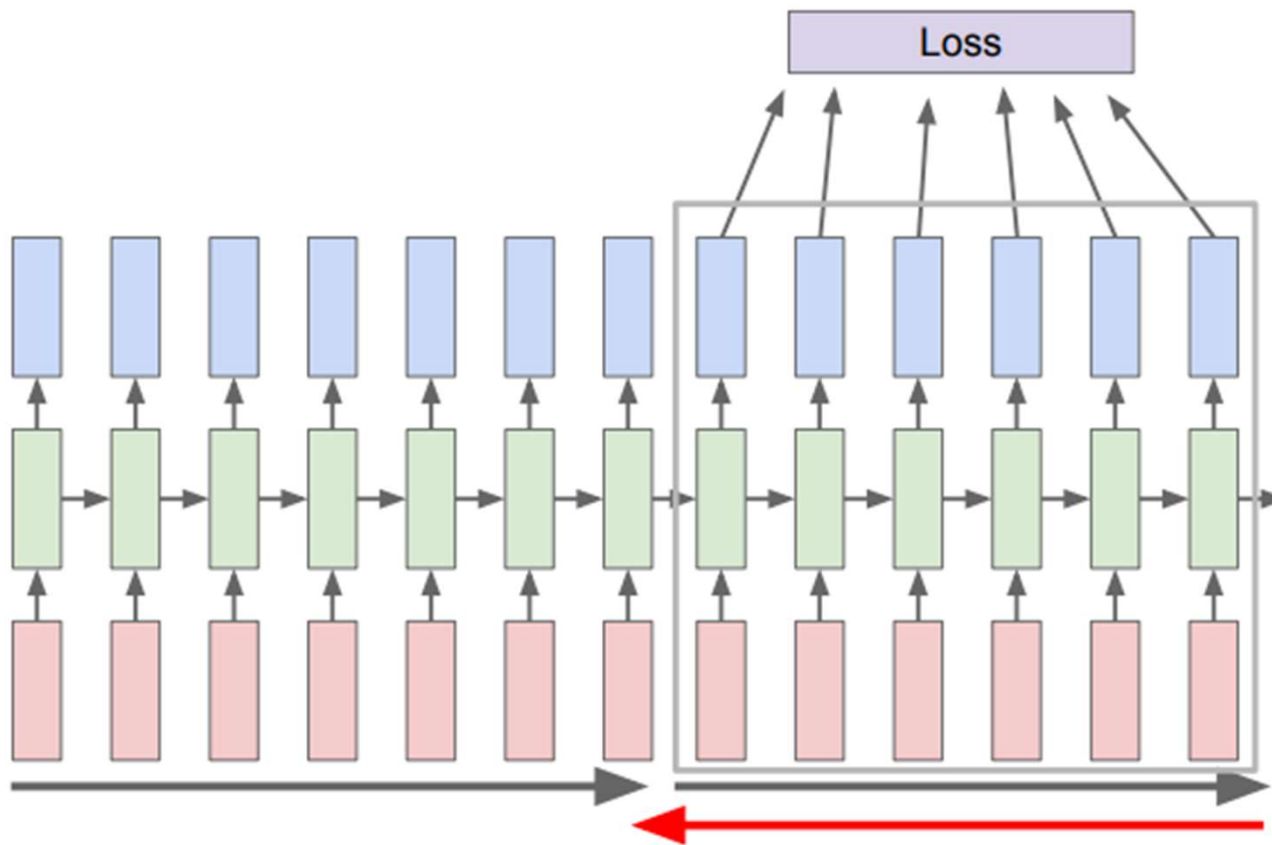


# Truncated Backpropagation through time




Run forward and backward through chunks of the sequence instead of whole sequence

# Truncated Backpropagation through time



Carry hidden states forward in time forever, but only backpropagate for some smaller number of steps


# Standard RNN Architecture



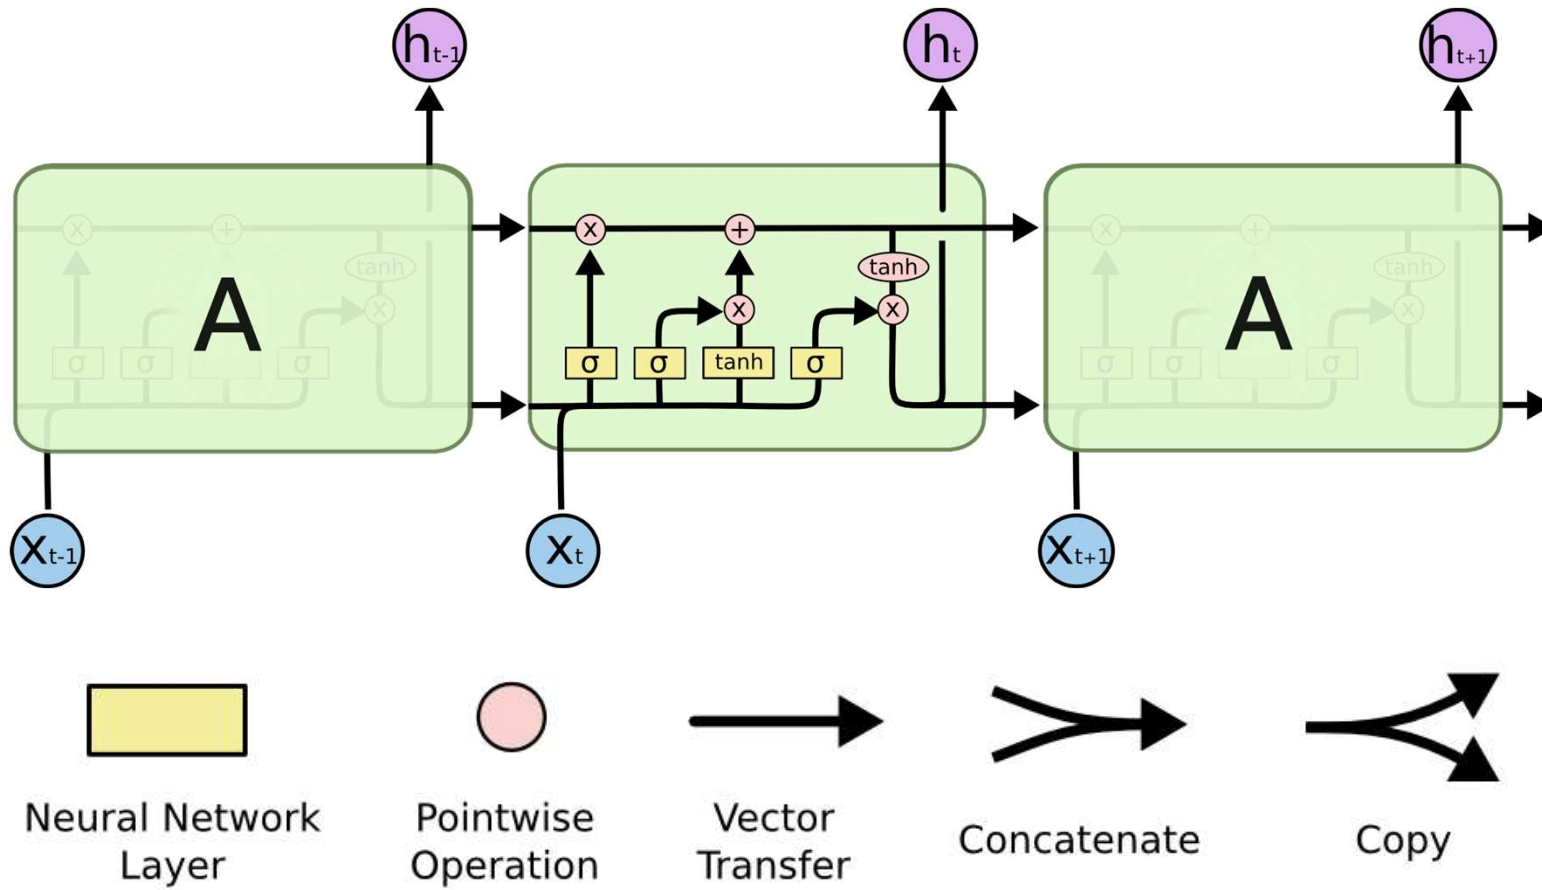
**The repeating module in a standard RNN contains a single layer.**

# Long Short-Term Memory

- LSTM networks, add additional gating units in each memory cell.
  - Forget gate
  - Input gate
  - Output gate
- Prevents vanishing/exploding gradient problem and allows network to retain state information over longer periods of time.




# LSTM Network Architecture



# Cell State

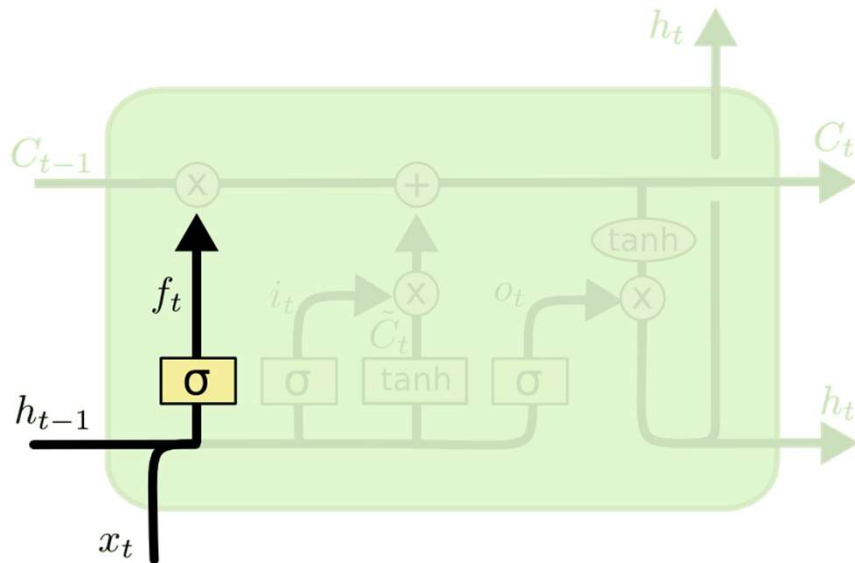
- Maintains a vector  $C_t$  that is the same dimensionality as the hidden state,  $h_t$
- Information can be added or deleted from this state vector via the forget and input gates.





# Forget Gate


- Forget gate computes a 0-1 value using a logistic sigmoid output function from the input,  $x_t$ , and the current hidden state,  $h_{t-1}$ :
- Multiplicatively combined with cell state, "forgetting" information where the gate outputs something close to 0.



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

# Input Gate

- First, determine which entries in the cell state to update by computing 0-1 sigmoid output.
- Then determine what amount to add/subtract from these entries by computing a tanh output (valued  $-1$  to  $1$ ) function of the input and hidden state.




$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

# Updating the Cell State


- Cell state is updated by using component-wise vector multiply to "forget" and vector addition to "input" new information.



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

# Output Gate

- Hidden state is updated based on a "filtered" version of the cell state, scaled to  $-1$  to  $1$  using  $\tanh$ .
- Output gate computes a sigmoid function of the input and current hidden state to determine which elements of the cell state to "output".




$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

$$\begin{pmatrix} f_t \\ i_t \\ o_t \\ \tilde{C}_t \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W_g \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$C_t = f_t \odot c_{t-1} + i_t \odot \tilde{C}_t$$

$$h_t = o_t \odot \tanh C_t$$



LSTM [Hochreiter et al., 1997]

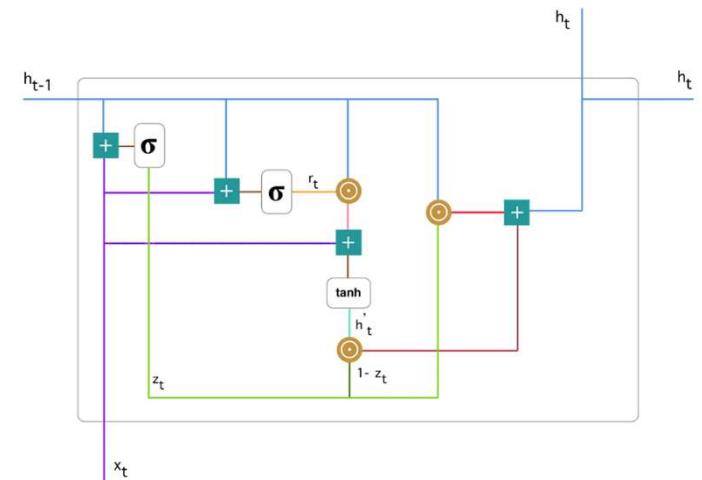
- **Forget gate ( $f_t$ ):** Defines how much of the previous state you want to let through.
- **Input gate ( $i_t$ ):** Defines how much of the newly computed state for the current input you want to let through.
- **Output gate ( $o_t$ ):** Defines how much of the internal state you want to expose to the external network.
- $\tilde{C}_t$ : “candidate” hidden state that is computed based on the current input and the previous hidden state.
- $C_t$ : the internal memory of the unit. Intuitively it is a combination of how we want to combine previous memory and the new input.
- Given the memory  $C_t$  we finally compute the output **hidden state  $h_t$**  by multiplying the memory with the output gate.

# Do LSTMs solve the vanishing gradient problem?

- The LSTM architecture makes it easier for the RNN to preserve information over many timesteps.
  - e.g., if the  $f = 1$  and the  $i = 0$ , then the information of that cell is preserved indefinitely.
  - By contrast, it's harder for vanilla RNN to learn a recurrent weight matrix  $W$  that preserves info in hidden state
- LSTM doesn't guarantee that there is no vanishing/exploding gradient, but it does provide an easier way for the model to learn long-distance dependencies.

# Gated Recurrent Unit (GRU)

1. Update gate:  $z_t = \sigma(W^{(z)}x_t + U^{(z)}h_{t-1})$
2. Reset gate:  $r_t = \sigma(W^{(r)}x_t + U^{(r)}h_{t-1})$
3. New memory content:  $h'_t = \tanh(Wx_t + r_t \odot Uh_{t-1})$
4. Final memory:  $h_t = z_t \odot h_{t-1} + (1 - z_t) \odot h'_t$



# RNN Summary

- Can process any length input. Computation for step  $t$  can use information from many steps back.
- Vanilla RNNs are simple but don't work very well - Common to use LSTM or GRU: their additive interactions improve gradient flow.
- Model size doesn't increase for longer input - Same weights applied on every timestep, so there is symmetry in how inputs are processed.
- LSTMs, better at capturing long-term dependencies compared to vanilla RNNs, may still struggle with very long sequences or maintaining context over extended periods.
- Computationally Intensive, Difficult in Parallelization, Limited Interpretability.
- Architectures like Transformers with their self-attention mechanisms have addressed these.



# Further Readings

- <https://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- [https://cs231n.stanford.edu/slides/2023/lecture\\_8.pdf](https://cs231n.stanford.edu/slides/2023/lecture_8.pdf)
- [https://cs231n.stanford.edu/slides/2020/lecture\\_10.pdf](https://cs231n.stanford.edu/slides/2020/lecture_10.pdf)